Topic：Understanding Linear Regression: Impact of Experience on Salary with Performance Evaluation and Critical Analysis

Student Name: Nimra Khizar

Course Title: Machine Learning and Neutral Networking

Introduction:

Machine Learning is a field of computer science that enables systems to learn patterns from data and make predictions. One of the most fundamental and widely used algorithms in supervised learning is **Linear Regression**. This tutorial focuses on understanding how years of experience affect salary using Linear Regression. It also explores model performance, assumptions, and limitations, making it useful for beginners who want to apply this technique in real-world scenarios.

In addition, this tutorial not only demonstrates how Linear Regression works but also evaluates its reliability using validation techniques and error analysis, making it more applicable to real-world machine learning problems.

Objective:

The objectives of this tutorial are:

- To understand the relationship between experience and salary
- To demonstrate how Linear Regression works mathematically and practically
- To evaluate model performance using MSE and R^2 score
- To critically analyze strengths and limitations of the model

Dataset Description:

This tutorial uses a dataset containing two variables:

Variable	Description	Type
Years Experience	Independent variable	Numeric
Salary	Dependent variable	Numeric

Why this dataset?

This dataset is simple and ideal for beginners. It clearly shows a linear relationship, making it perfect for demonstrating Linear Regression concepts. The dataset is **synthetic but realistic,** commonly used in ML tutorials for educational purposes.

Theory of Linear Regression:

Linear Regression models the relationship between variables using:

$$y = mx + b$$

Where:

y = predicted salary

x = years of experience

m = slope (rate of increase)

b = intercept

How it works?

The model finds the best-fit line by minimizing the error between predicted and actual values using **Least Squares Method.** The Least Squares Method minimizes the sum of squared residuals (the difference between actual and predicted values). This ensures that the regression line is the best possible fit for the given data.

Mathematically, the objective is to minimize:

$\sum(y_i - \hat{y}_i)^2$

This optimization makes Linear Regression highly efficient for prediction tasks involving continuous variables.

Assumptions of Linear Regression:

For accurate results, Linear Regression assumes:

Linearity: Relationship must be linear

Independence: Observations are independent

Homoscedasticity: Constant variance of errors

No multicollinearity: Inputs not highly correlated

Normal distribution of errors

Note: If these assumptions are violated, model performance decreases.

Overfitting and Underfitting:

Overfitting occurs when a model learns not only the underlying pattern but also the noise in the data, leading to poor performance on unseen data.

Underfitting occurs when the model is too simple to capture the underlying trend.

In this tutorial, Linear Regression avoids overfitting due to its simplicity. However, it may underfit if the real-world relationship is non-linear.

Methodology:

Steps followed:

1. Dataset creation
2. Data visualization
3. Model training using Linear Regression
4. Prediction
5. Model evaluation (MSE & R^2 score)
6. Train-test split for model validation

Train-Test Split:

To ensure that the model generalizes well to unseen data, the dataset is divided into training and testing sets. The model is trained on the training data and evaluated on the testing data to avoid overfitting. The dataset is divided into training and testing sets using an 80-20 ratio.

Implementation (Python Code):

```
# Import Libraries

import pandas as pd

import matplotlib.pyplot as plt

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, r2_score

from sklearn.model_selection import train_test_split

# Create & Load Dataset

data = {

   "YearsExperience": [1.1,1.5,2.0,2.2,3.0,3.2,4.0,4.5,5.0,5.5,6.0,6.8,7.0,8.0,9.0,10.0],

   "Salary":
[39343,46205,43525,39891,56642,60150,55794,61111,66029,67938,81363,91738,93940,101302,105582,109431]
```

```
}
df = pd.DataFrame(data)
df.head()
# Features
X = df[['YearsExperience']]
y = df['Salary']
# Train- Test Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
# Model Training
model = LinearRegression()
model.fit(X_train, y_train)
# Prediction
y_pred = model.predict(X_test)
# Evaluation
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("MSE:", mse)
print("R² Score:", r2)
# Data Visualization & Graphs
# 1. Linear Regression (Test Data)
plt.scatter(X_test, y_test)
plt.plot(X_test, y_pred)
plt.xlabel("Years of Experience")
plt.ylabel("Salary")
```

plt.title("Linear Regression (Test Data)")

plt.show()

 # 2. Residual Plot

residuals = y_test - y_pred

plt.scatter(y_pred, residuals)

plt.axhline(y=0)

plt.xlabel("Predicted Values")

plt.ylabel("Residuals")

plt.title("Residual Plot")

The dataset was split into training and testing sets (80% training, 20% testing) to evaluate model performance on unseen data. This helps ensure that the model generalizes well and avoids overfitting. The dataset is split into training and testing sets using an 80-20 ratio. The model is trained on training data and predictions are made on test data.

Data Visualization:

1. **Scatter Plot for training and testing data + regression Line (Experience vs Salary):**

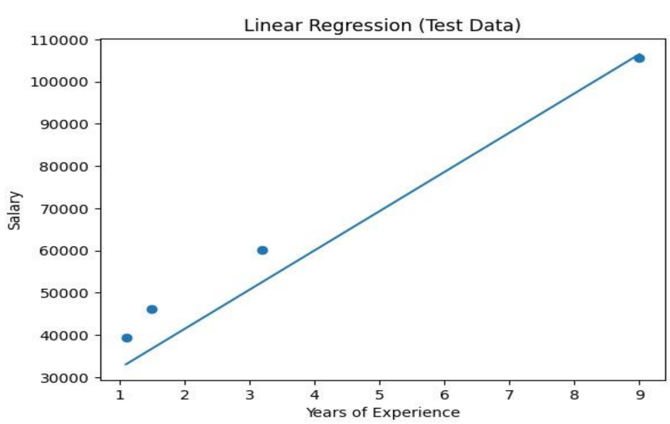

Figure 1: Linear relationship between experience and salary with regression line.

Observations:

- The plot shows a **linear regression** line modeling the relationship between **Years of Experience** (X-axis) and **Salary** (Y-axis).
- The regression line indicates a **positive linear correlation:** salary increases as experience increases.
- Points are close to the line, suggesting the linear model is a good fit for the data.

2. Residual Plot

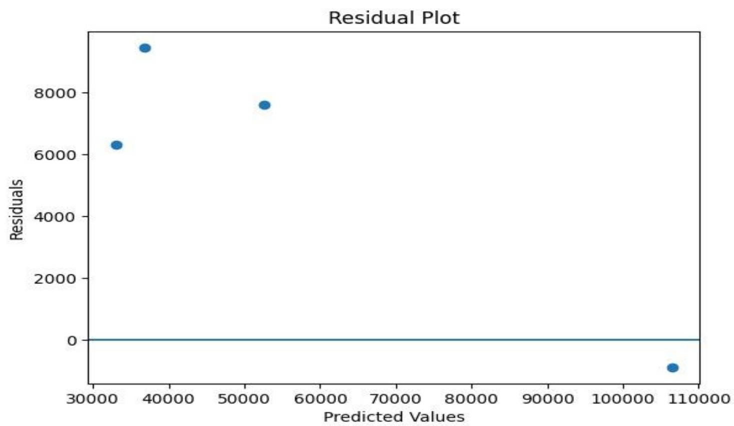

Figure 2: Relationship between predicted values and

Residuals with regression line.

Observations:

- Residuals represent the **difference** between actual salary values and the model's predicted values.
- Residuals are scattered randomly around the zero line, showing the model's predictions are unbiased and the errors are normally distributed.
- The absence of a pattern in residuals confirms the **linear regression** assumption is valid for this data.
- Residual plots help evaluate whether the assumptions of linear regression are satisfied.
- Random distribution of residuals around zero indicates a good model fit.

Results and Discussion:

R^2 Score ≈ High (~0.95+): Model explains most of the variance

MSE = Low: Predictions are close to actual values

Interpretation: This indicates that Linear Regression performs very well on this dataset.

However, since the dataset is small and synthetic, the high R^2 score may not generalize well to real-world data. Therefore, model validation using unseen data is essential. Although the model shows high accuracy, the dataset is small and synthetic. Therefore, the results may not generalize well to real-world scenarios. Testing on unseen data improves reliability. The model performs well on this dataset; however, performance may decrease on real-world data due to noise and non-linearity.

Critical Analysis:

Strengths:

- Simple and easy to implement
- Highly interpretable
- Works well with linear data

Limitations:

- Cannot handle non-linear relationships
- Sensitive to outliers
- Assumes perfect linearity

Additionally, Linear Regression assumes a strictly linear relationship, which is rarely the case in real-world scenarios. The model is also sensitive to outliers, which can significantly distort predictions.

Real-Life Applications:

- Salary prediction
- House price estimation
- Sales forecasting
- Business trend analysis

Personal Reflection:

While working on this tutorial, I learned that Linear Regression is a powerful yet simple tool for prediction tasks. However, I also realized that it is not suitable for complex datasets where relationships are non-linear. This helped me understand the importance of choosing the right model based on data. If I had more time, I would extend this project by applying Polynomial Regression and comparing its performance with Linear Regression to better understand non-linear patterns.

Ethical Considerations:

Machine learning models can introduce bias if the dataset is not fair. In salary prediction, bias may lead to discrimination. It is important to ensure fairness and transparency in model usage. In salary prediction, biased data (e.g., gender or regional bias) may lead to unfair outcomes. Therefore, fairness and transparency must be ensured when deploying machine learning models.

Conclusion:

Linear Regression effectively models the relationship between experience and salary. It provides accurate predictions when assumptions are satisfied. This tutorial demonstrates how beginners can apply machine learning techniques to real-world problems.

References:

- Montgomery, D.C. (2021). Introduction to Linear Regression Analysis.
- Scikit-learn (2024). Linear Models Documentation. Available at: https://scikit-learn.org/stable/modules/linear_model.html
- GeeksforGeeks (2023). Linear Regression in Python. Available at: https://www.geeksforgeeks.org/linear-regression-python-implementation/
- Towards Data Science (2023). Understanding Linear Regression. Available at: https://towardsdatascience.com/
- Bishop, C. (2006). Pattern Recognition and Machine Learning. Springer.
- Hastie, T., Tibshirani, R., & Friedman, J. (2009). The Elements of Statistical Learning.